

Autonomous Learning

Assignment 2

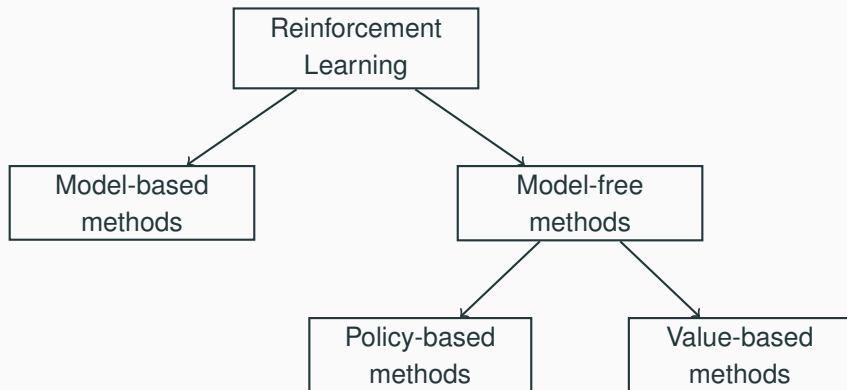
Simon Reichhuber

June 15, 2020

University of Kiel, Summer Term 2019

1. Reinforcement Learning
2. Artificial Neural Networks
3. Deep Q-Network
4. Bibliography

Reinforcement Learning



- Agent learns a model of the environment
- Action a_1 in state $s_1 \rightarrow$ state s_2 and reward r_2
 \Rightarrow improvement of the estimates of $T(s_2|s_1, a_1)$ and $R(s_1, a_1)$
- As soon as the model is sufficient \rightarrow policy


Example: Value Iteration and Policy Extraction

Value-based methods

... are based on *temporal difference learning*: They learn the value function V^π or V^* or the Q function Q^π or Q^*

Policy-based methods

... directly learn the optimal policy π^* (or try to approximate the optimal policy in case the real optimal policy is not reachable)



- Agents learn the value of the policy in use to make decisions
- The estimated value function is updated using the results of actions chosen by policy π

Example: **SARSA**

- The estimated value function can be updated by **hypothetical actions**: actions that are not explicitly explored

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a' \in A} Q(s', a') - Q(s, a)]$$

- The agent learns strategies that he not necessarily explored during training.

Example: **Q-Learning**

- On-policy
- An episode consists of an alternating sequence of states and state-action-pairs:



- Learning the state-action-pairs:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Quintuple: $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}) \rightarrow \text{SARSA}$

SARSA learns $Q^\pi(S_t, A_t)$

$$Q^\pi(S_t, A_t) \leftarrow Q^\pi(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q^\pi(S_{t+1}, A_{t+1}) - Q^\pi(S_t, A_t) \right]$$

Q-Learning learns $Q^*(S_t, A_t)$

$$Q^*(S_t, A_t) \leftarrow Q^*(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q^*(S_{t+1}, a) - Q^*(S_t, A_t) \right]$$


Both converge to Q^π resp. Q^* if enough samples for each state-action-pair are given.


- The agent starts in S
- The target is blue final state (positive reward)
- But there are cliffs on the way (negative reward and restart!)

S															


Q-Learning \rightarrow highest action value

BUT: Exploration \rightarrow random action





- optimal (fastest) route
- not save



- longer route
- save
- good online performance

Algorithm 1: SARSA

Initialize $Q(s, a)$ arbitrarily and the $Q(\text{terminal}, \cdot) = 0$; **repeat**

 Initialize s ;

 Choose a from s using policy derived from Q (e.g. ϵ -greedy);

repeat

 Take action a , observe r, s' ;

 Choose a' from s' using policy derived from Q (e.g. ϵ -greedy);

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$;


$s \leftarrow s'$;


$a \leftarrow a'$;

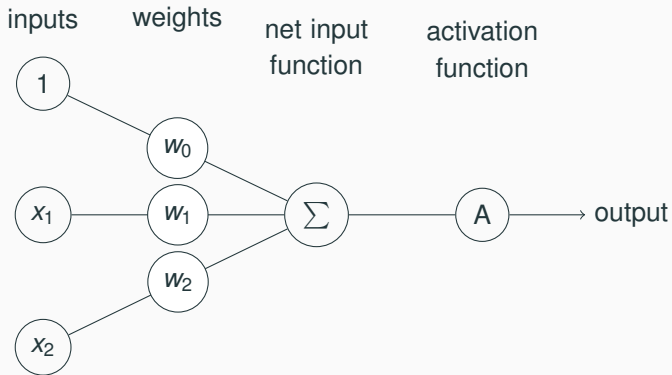
until s is terminal;

until Q is converged;

Artificial Neural Networks








- A node is fired when the inputs meet certain requirements
- The **input function** sums up all inputs ($input \times weight$)
- The **activation function** decides if and how intense a signal is sent

- The optimisation function adapts the weights according to the error they produce.
- Gradient denotes the relationship between a single weight and the error of the whole network.
- Slow adaption of many weights → Which input has what significance?
- Chain Rule:


$$\frac{dError}{dweight} = \frac{dError}{dactivation} * \frac{dactivation}{dweight}$$

- Learning: Adaption of the weights until the error cannot be minimised any further.

linear:



ReLu (Rectifier Linear Unit):



- The error for trainable weights θ is generally defined as the difference between the predicted and the actual output.


$$J(\theta) = p - \hat{p}$$

- The function for calculating the error is called Loss Function J (or cost function).
- Different Loss Functions \rightarrow different errors
- Frequently used Loss Function:

mean square error (MSE)

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n (p_i - \hat{p}_i)^2$$

- $J(\theta) = J(w)$ = function of the internal parameters (weights and bias)
- The error is passed through the layers from the back to front.



- A high-level neural networks API.
- Different libraries can work in the background (e. g. Tensorflow, CNTK, Theano)
- www.keras.io

```
1 from keras.models import Sequential
2 from keras.layers import Dense
3 from keras.optimizers import RMSprop
4
5 model = Sequential()
6
7 # input and first hidden layer
8 model.add(Dense(units=128, activation='relu', input_dim=2))
9
10 # second hidden layer
11 model.add(Dense(units=256, activation='relu'))
12
13 # output layer
14 model.add(Dense(units=8, activation='linear'))
15
16 model.compile(loss='mse',
17               optimizer=RMSprop(lr=0.00025))
18
19 # train network
20 model.fit(x_train, y_train, batch_size=32)
21
22 # predict on trained network
23 prediction = model.predict(x_test)
```

Deep Q-Network

- Large state space and/or action space \rightarrow very large Q-table
 \Rightarrow approximation of the Q-table using a neural net
- A neural net can generalize its knowledge of visited states for non-visited states
- Abstraction of patterns and understanding actions on the basis of already seen patterns.


- State space represented as vector
- Loss Function:

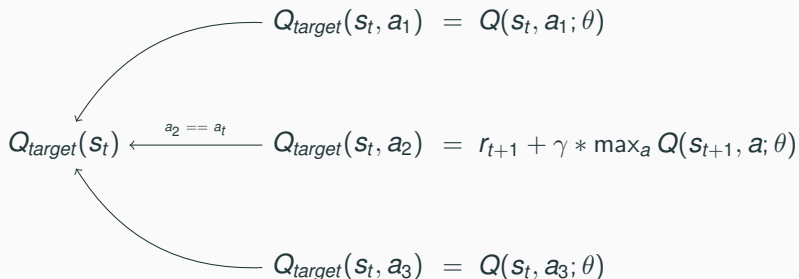
$$J(\theta) = \sum (Q - Q_{target})^2$$

- Approximating function:

$$Q(s_t, a_t; \theta) \leftarrow Q(s_t, a_t; \theta) + \alpha \left[\underbrace{r_t + \gamma \max_a Q(s_{t+1}, a; \theta)}_{target} - Q(s_t, a_t; \theta) \right]$$

$TD-error$




$$Q_{target}(s_t) \leftarrow \begin{array}{l} Q_{target}(s_t, a_1) = Q(s_t, a_1; \theta) \\ \xleftarrow{a_2 == a_t} Q_{target}(s_t, a_2) = r_{t+1} + \gamma * \max_a Q(s_{t+1}, a; \theta) \\ Q_{target}(s_t, a_3) = Q(s_t, a_3; \theta) \end{array}$$

- Non-linear approximation function (ANN) \rightarrow unstable learning
- Reasons:
 - Correlation between some observations
 - Correlation between action and target values
 - Small adaptations can lead to significant changes of the policy and subsequently also the distribution of the data.

- Remember the last N transitions ($s_t, a_t, r_{t+1}, s_{t+1}, done$)
- Instead of learning from just the last transition: For each step draw a random minibatch (of size 32) from the experience replay.
- Q-Learning Updates based on this minibatch
- FiFo



- This removes correlations between individual observations and smoothes changes in the data distribution.
- Transitions are used more often for learning \rightarrow data efficiency

- In every step the values of the Q-network shift
⇒ feedback loops between target values and predicted q-values
- Target network: A second neural network used during training
- Calculates the target Q-values for the Loss Function
- Is periodically updated (every C steps)



- Reduces the correlation between action and target values

- Defining a minimal and maximal error:

$$\left[r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right] \in [-1, 1]$$

⇒ more stable learning

θ = weights of the Q-network

θ^- = weights of the target network

Loss Function:

$$J(\theta) = \sum \left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2$$

Algorithm 2: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

for $episode=1, M$ **do**

for $t=1, T$ **do**

 With probability ϵ select random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(s_t, a; \theta)$

 Execute action a_t and observe reward r_{t+1} and state s_{t+1}

 Store transition $(s_t, a_t, r_{t+1}, s_{t+1}, done)$ in D

 Sample random minibatch of transitions $(s_j, a_j, r_{j+1}, s_{j+1}, done)$ from D

 Set $y_j = \begin{cases} r_j & \text{if } done \\ r_j + \gamma \max_{a'} \hat{Q}(s_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 perform a gradient descent step on $(y_j - Q(s_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

Bibliography

- <https://www.quora.com/What-is-the-differences-between-artificial-neural-network-computer-science-and-biological-neural-network>
- <https://www.quora.com/In-neural-networks-how-important-is-back-propagation-What-is-its-significance>